

Lecture 2: xv6 introduction and x86 background



References: Appendix A of xv6 book

Chapters 3,4 from Programming from the
Ground Up

xv6

- ❖ xv6 is a simple OS for easy teaching of OS concepts
 - Two versions, one for x86 hardware and one for RISC-V hardware
 - This series of lectures based on **x86 version**
 - <https://github.com/mit-pdos/xv6-public>
- ❖ This lecture: overview of x86 hardware and other background needed to understand xv6 code

Understanding xv6: background

- ❖ OS enables processes stored in memory to run on CPU
 - Process code/data in main memory
 - CPU fetches, decodes, executes instructions in program code
 - Process data fetched from memory to CPU registers for faster access during instruction execution
 - Recently fetched code/data stored in CPU caches for future access (memory access is very slow compared to CPU)
- ❖ What we will cover in this lecture
 - Common x86 registers
 - Common x86 instructions
 - How stack is used during function calls (C calling convention)

How a Process is stored in main Memory?

- ❖ Process memory image consists of
 - Compiled code (CPU instructions)
 - Global/static variables (memory allocated at compile time)
 - Heap (dynamic memory allocation via, e.g., malloc) that grows on demand
 - Stack (temporary storage during function calls, e.g., local variables) that usually grows “up” towards lower addresses
 - Other things like shared libraries
- ❖ Every instruction/data has an address, used by CPU to fetch/store
 - Virtual addresses (managed by OS)
- ❖ Example: can you understand what is happening with variable “iptr”?

```
int *iptr = malloc(sizeof(int))
```

x86 registers: examples

- ❖ Small space for data storage within CPU
- ❖ General purpose registers: store data during computations (eax, ebx, ecx, edx, esi, edi)
- ❖ Pointers to stack locations: base of stack (ebp) and top of stack (esp)
- ❖ Program counter or instruction pointer (eip): next instruction to execute
- ❖ Control registers: hold control information or metadata of a process (e.g., cr3 has pointer to page table)
- ❖ Segment registers (cs, ds, es, fs, gs, ss): information about segments (related to memory of process)

x86 instructions: examples

◊ Move Instruction (Load/store)

◊ AT&T syntax i.e.

◦ *mov %eax, %ebx*

mov src, dst

(copy contents of eax to ebx)

◦ *mov (%eax), %ebx*

(copy contents at the address in eax into ebx)

◦ *mov 4(%eax), %ebx*

(copy contents stored at offset of 4 bytes from
address stored at eax into ebx)

at
address
of

x86 instructions: examples

- ❖ Push/pop on stack: changes esp
 - *push %eax* (push contents of eax onto stack, update esp)
 - ~~*pop %eax*~~ (pop top of stack onto eax, update esp)
- ❖ *jmp* sets eip to specified address
- ❖ *call* to invoke a function, *ret* to return from a function
- ❖ Other variants
 - (*movw*, *pushl*) for different register sizes

Privilege levels

- ◊ x86 CPUs have multiple privilege levels
 - Four “rings” (0 to 3)
 - Ring 0 has highest privilege, runs OS code
 - Ring 3 has lowest privilege, runs user code
- Two types of instructions:
 - privileged and unprivileged

0 → highest

3 → lowest

Privileged Instructions

- ◊ Privileged instructions can be executed by CPU only when running at the highest privilege level (ring 0)
 - For example, writing into cr3 register (setting page table) is privileged instruction, only OS should do it, because we do not want a user manipulating memory of another process
 - Another example: instructions to access I/O devices

Unprivileged instructions

- Unprivileged instructions can be run at lower privilege levels
 - For example, user code running at lower privilege can store a value into a general purpose register
- When user requires OS services (e.g., system call), CPU moves to higher privilege level and executes OS code that contains privileged instructions
 - User code cannot invoke privileged instructions directly

Function calls

- ◊ Local variables, arguments stored on stack for duration of function call
- ◊ What happens in a function call?
 - Push function arguments on stack
 - call fn (instruction pushes return address on stack, jumps to function)
 - Allocate local variables on stack
 - Run function code
 - ret (instruction pops return address, eip goes back to old value)

Register Savings

- ❖ Register savings is done to prevent overwriting
 - Register can get clobbered during function call
 - Two methods
 - ❖ Caller saved : registers saved on stack by caller before invoking the function (caller save registers). Function code (callee) can freely change them, caller restores them later on.
 - Callee saved: Some registers saved by callee function and restored after function ends (callee save registers). Caller expects them to have same value on return.
 - Return value stored in eax register by callee (one of caller save registers)
- ❖ All of this is automatically done by C compiler (C calling convention)

caller saved,
callee saved, (same
register saved by
callee)

Timeline: Function calls & Stack

- ◊ Timeline of a function call (note: stack grows from "up" from higher to lower addresses)

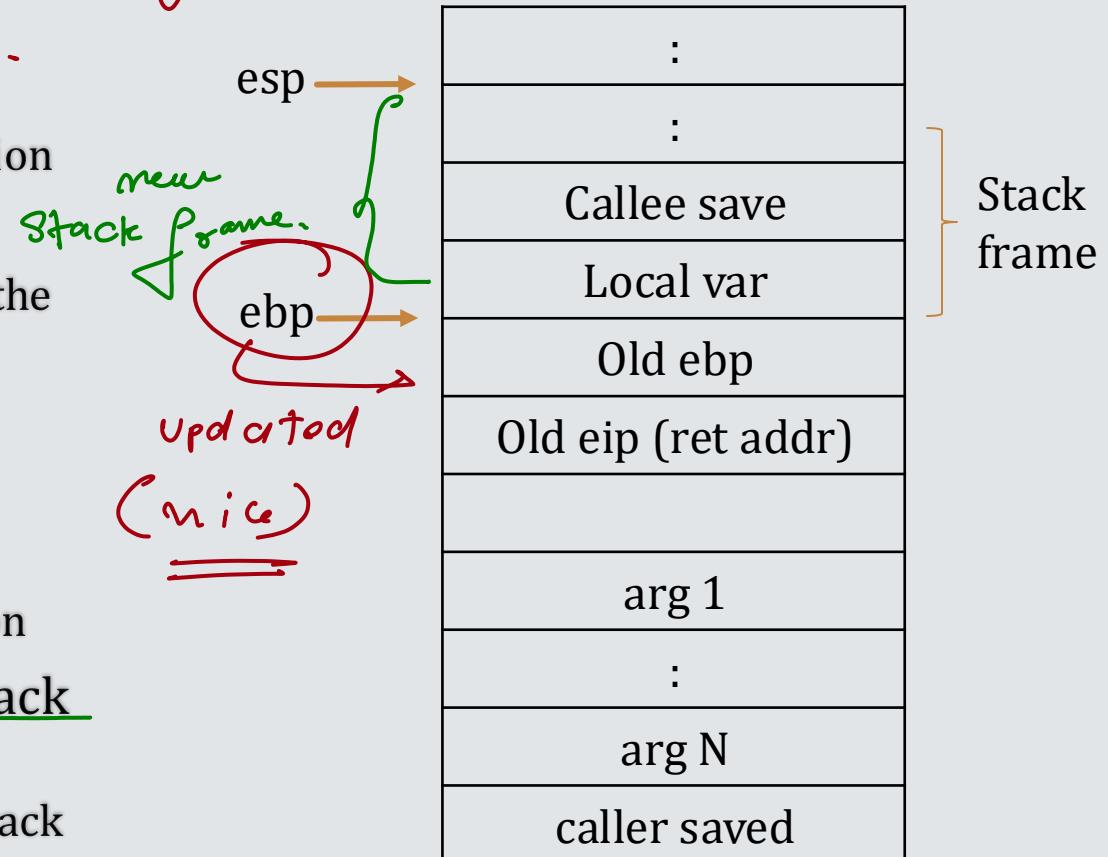
- Push caller save registers (eax, ecx, edx)
- Push arguments in reverse order
- Return address (old eip) pushed on stack by the call instruction
- Push old ebp on stack
- Set ebp to current top of stack (base of new "stack frame" of the function)
- Push local variables and callee save registers (ebx, esi, edi)
- Execute function code
- Pop stack frame and restore old ebp
- Return address popped and eip restored by the ret instruction

- ◊ Stack pointers: ebp stores address of base of current stack frame and esp stores address of current top of stack

- Function arguments are accessible from looking under the stack base pointer

(using offset)

Stack grows from Higher to Lower address



C vs. assembly for OS code

❖ Why all this x86 background?

- Most of xv6 is in C, and assembly code (including all the stack manipulations for function calls) is automatically generated by compiler.
- However, small parts are in assembly language. Why?
- Sometimes, OS needs more control over what needs to be done (for example, the logic of switching from stack of one process to stack of another cannot be written in a high-level language)

❖ Basic understanding of x86 assembly language is required to follow some nuances of xv6 code



Thank You